

Neeti Capstone project 2 (Healthcare)

April 8, 2022

```
[1]: import pandas as pd
```

```
[2]: data=pd.read_csv("health care diabetes.csv")
df=pd.read_csv("healthcare appointment data.csv")
df_train=pd.read_csv("train.csv")
```

```
[3]: data
```

```
[3]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|-----|-------------|---------|---------------|---------------|---------|------|---|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | |
| .. | ... | ... | ... | ... | ... | ... | |
| 763 | 10 | 101 | 76 | 48 | 180 | 32.9 | |
| 764 | 2 | 122 | 70 | 27 | 0 | 36.8 | |
| 765 | 5 | 121 | 72 | 23 | 112 | 26.2 | |
| 766 | 1 | 126 | 60 | 0 | 0 | 30.1 | |
| 767 | 1 | 93 | 70 | 31 | 0 | 30.4 | |

| | DiabetesPedigreeFunction | Age | Outcome |
|-----|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |
| 3 | 0.167 | 21 | 0 |
| 4 | 2.288 | 33 | 1 |
| .. | ... | ... | ... |
| 763 | 0.171 | 63 | 0 |
| 764 | 0.340 | 27 | 0 |
| 765 | 0.245 | 30 | 0 |
| 766 | 0.349 | 47 | 1 |
| 767 | 0.315 | 23 | 0 |

[768 rows x 9 columns]

```
[4]: type(data)
```

```
[4]: pandas.core.frame.DataFrame
```

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome               768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[6]: data.describe()
```

```
[6]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin \ |
|-------|-------------|------------|---------------|---------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 |

| | BMI | DiabetesPedigreeFunction | Age | Outcome |
|-------|------------|--------------------------|------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 31.992578 | 0.471876 | 33.240885 | 0.348958 |
| std | 7.884160 | 0.331329 | 11.760232 | 0.476951 |
| min | 0.000000 | 0.078000 | 21.000000 | 0.000000 |
| 25% | 27.300000 | 0.243750 | 24.000000 | 0.000000 |
| 50% | 32.000000 | 0.372500 | 29.000000 | 0.000000 |
| 75% | 36.600000 | 0.626250 | 41.000000 | 1.000000 |
| max | 67.100000 | 2.420000 | 81.000000 | 1.000000 |

```
[7]: data.isnull()
```

```
[7]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI \ |
|---|-------------|---------|---------------|---------------|---------|-------|
| 0 | False | False | False | False | False | False |

| | | | | | | | |
|-----|-------|-------|-------|-------|-------|-------|-------|
| 1 | False | False | False | False | False | False | False |
| 2 | False | False | False | False | False | False | False |
| 3 | False | False | False | False | False | False | False |
| 4 | False | False | False | False | False | False | False |
| .. | ... | ... | ... | ... | ... | ... | ... |
| 763 | False | False | False | False | False | False | False |
| 764 | False | False | False | False | False | False | False |
| 765 | False | False | False | False | False | False | False |
| 766 | False | False | False | False | False | False | False |
| 767 | False | False | False | False | False | False | False |

| | DiabetesPedigreeFunction | Age | Outcome |
|-----|--------------------------|-------|---------|
| 0 | False | False | False |
| 1 | False | False | False |
| 2 | False | False | False |
| 3 | False | False | False |
| 4 | False | False | False |
| .. | ... | ... | ... |
| 763 | False | False | False |
| 764 | False | False | False |
| 765 | False | False | False |
| 766 | False | False | False |
| 767 | False | False | False |

[768 rows x 9 columns]

```
[8]: miss_cols = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
, 'BMI']
```


```
[9]: #Importing libraries


import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```


```
[10]: for col in miss_cols:
    plt.figure(figsize = (5, 5))
    plt.title(col)
    sns.distplot(data[col], kde = False, color = 'red')
```


C:\Users\91820\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).


```
warnings.warn(msg, FutureWarning)
```



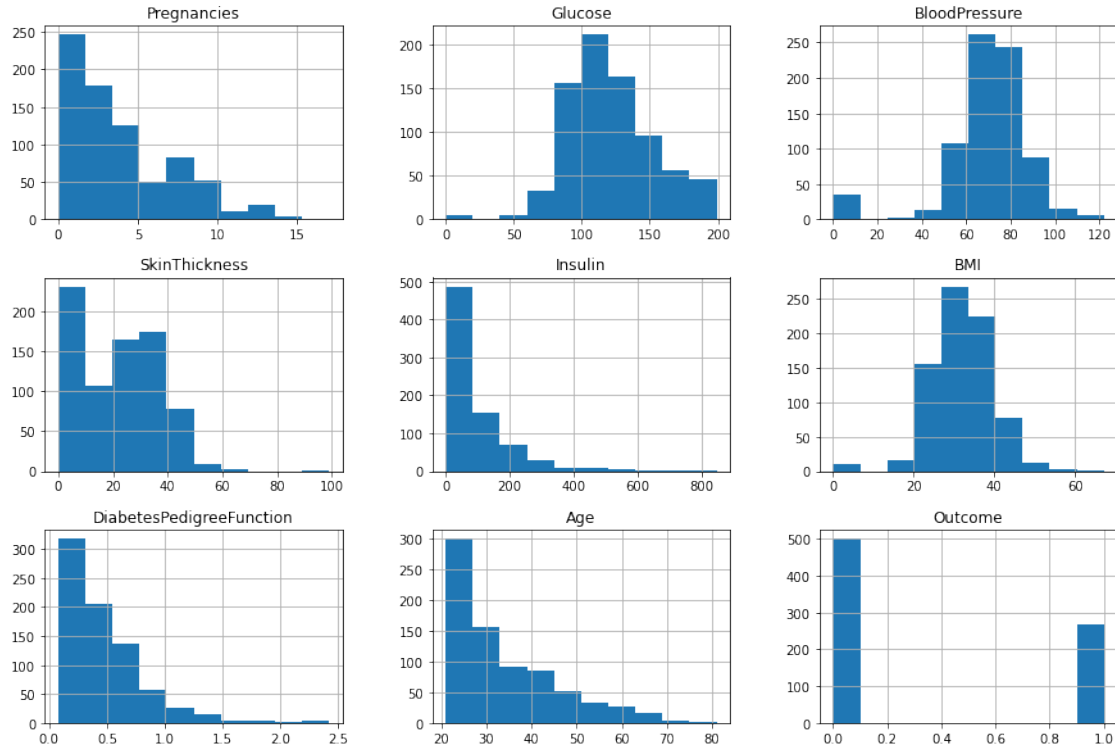









```
[11]: data.hist(figsize=(15,10))  
plt.show()
```





```
[12]: for column in data.columns:
      sns.boxplot(data[column])
      plt.show()
```


C:\Users\91820\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(




```
C:\Users\91820\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```




```
C:\Users\91820\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```




```
C:\Users\91820\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```




```
C:\Users\91820\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```




```
C:\Users\91820\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```




```
C:\Users\91820\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```



```
C:\Users\91820\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```

```
C:\Users\91820\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```



```
[ ]:
```

```
[13]: data[['Glucose', 'BloodPressure', 'SkinThickness',
'Insulin', 'BMI']] = data[['Glucose', 'BloodPressure',
'SkinThickness', 'Insulin', 'BMI']].replace(0, np.nan)
```

```
[14]: data
```

```
[14]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|-----|-------------|---------|---------------|---------------|---------|------|---|
| 0 | 6 | 148.0 | 72.0 | 35.0 | NaN | 33.6 | |
| 1 | 1 | 85.0 | 66.0 | 29.0 | NaN | 26.6 | |
| 2 | 8 | 183.0 | 64.0 | NaN | NaN | 23.3 | |
| 3 | 1 | 89.0 | 66.0 | 23.0 | 94.0 | 28.1 | |
| 4 | 0 | 137.0 | 40.0 | 35.0 | 168.0 | 43.1 | |
| .. | ... | ... | ... | ... | ... | ... | |
| 763 | 10 | 101.0 | 76.0 | 48.0 | 180.0 | 32.9 | |
| 764 | 2 | 122.0 | 70.0 | 27.0 | NaN | 36.8 | |
| 765 | 5 | 121.0 | 72.0 | 23.0 | 112.0 | 26.2 | |
| 766 | 1 | 126.0 | 60.0 | NaN | NaN | 30.1 | |
| 767 | 1 | 93.0 | 70.0 | 31.0 | NaN | 30.4 | |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |

| | | | |
|-----|-------|-----|-----|
| 3 | 0.167 | 21 | 0 |
| 4 | 2.288 | 33 | 1 |
| .. | ... | ... | ... |
| 763 | 0.171 | 63 | 0 |
| 764 | 0.340 | 27 | 0 |
| 765 | 0.245 | 30 | 0 |
| 766 | 0.349 | 47 | 1 |
| 767 | 0.315 | 23 | 0 |

[768 rows x 9 columns]

```
[15]: pip install impyute
```

```
Requirement already satisfied: impyute in c:\users\91820\anaconda3\lib\site-
packages (0.0.8)
Requirement already satisfied: scikit-learn in
c:\users\91820\anaconda3\lib\site-packages (from impyute) (1.0.2)
Requirement already satisfied: numpy in c:\users\91820\anaconda3\lib\site-
packages (from impyute) (1.20.3)
Requirement already satisfied: scipy in c:\users\91820\anaconda3\lib\site-
packages (from impyute) (1.7.1)
Requirement already satisfied: joblib>=0.11 in
c:\users\91820\anaconda3\lib\site-packages (from scikit-learn->impyute) (1.1.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
c:\users\91820\anaconda3\lib\site-packages (from scikit-learn->impyute) (2.2.0)
Note: you may need to restart the kernel to use updated packages.
```

```
[16]: from impyute.imputation.cs import fast_knn
```

```
[17]: imputed_data = fast_knn(data.values, k = 30)
```

```
[18]: imputed_data
```

```
[18]: array([[ 6.    , 148.    , 72.    , ..., 0.627, 50.    , 1.    ],
 [ 1.    , 85.    , 66.    , ..., 0.351, 31.    , 0.    ],
 [ 8.    , 183.    , 64.    , ..., 0.672, 32.    , 1.    ],
 ...,
 [ 5.    , 121.    , 72.    , ..., 0.245, 30.    , 0.    ],
 [ 1.    , 126.    , 60.    , ..., 0.349, 47.    , 1.    ],
 [ 1.    , 93.    , 70.    , ..., 0.315, 23.    , 0.    ]])
```

```
[19]: imputed_data = pd.DataFrame(imputed_data)
```

```
[20]: imputed_data.head()
```

```
[20]:      0      1      2      3      4      5      6      7      8
0  6.0  148.0  72.0  35.000000  155.333764  33.6  0.627  50.0  1.0
```

```

1  1.0   85.0  66.0  29.000000  155.548223  26.6  0.351  31.0  0.0
2  8.0  183.0  64.0  29.367818  155.374337  23.3  0.672  32.0  1.0
3  1.0   89.0  66.0  23.000000   94.000000  28.1  0.167  21.0  0.0
4  0.0  137.0  40.0  35.000000  168.000000  43.1  2.288  33.0  1.0

```

```
[21]: data.head()
```

```

[21]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0           6    148.0           72.0           35.0      NaN  33.6
1           1     85.0           66.0           29.0      NaN  26.6
2           8    183.0           64.0           NaN      NaN  23.3
3           1     89.0           66.0           23.0     94.0  28.1
4           0    137.0           40.0           35.0    168.0  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627    50         1
1                0.351    31         0
2                0.672    32         1
3                0.167    21         0
4                2.288    33         1

```

```
[22]: data.columns
```

```

[22]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
        'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
        dtype='object')

```

```

[23]: imputed_data.columns = ['Pregnancies', 'Glucose', 'BloodPressure',
    ↪ 'SkinThickness', 'Insulin',
    'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']

```

```
[24]: imputed_data.head()
```

```

[24]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0           6.0    148.0           72.0    35.000000  155.333764  33.6
1           1.0     85.0           66.0    29.000000  155.548223  26.6
2           8.0    183.0           64.0    29.367818  155.374337  23.3
3           1.0     89.0           66.0    23.000000   94.000000  28.1
4           0.0    137.0           40.0    35.000000  168.000000  43.1


      DiabetesPedigreeFunction  Age  Outcome
0                0.627   50.0         1.0
1                0.351   31.0         0.0
2                0.672   32.0         1.0
3                0.167   21.0         0.0
4                2.288   33.0         1.0


```


```
[25]: for col in miss_cols:
      plt.figure(figsize = (5, 5))
      plt.title(col)
      sns.distplot(imputed_data[col], kde = False, color = 'red')
```

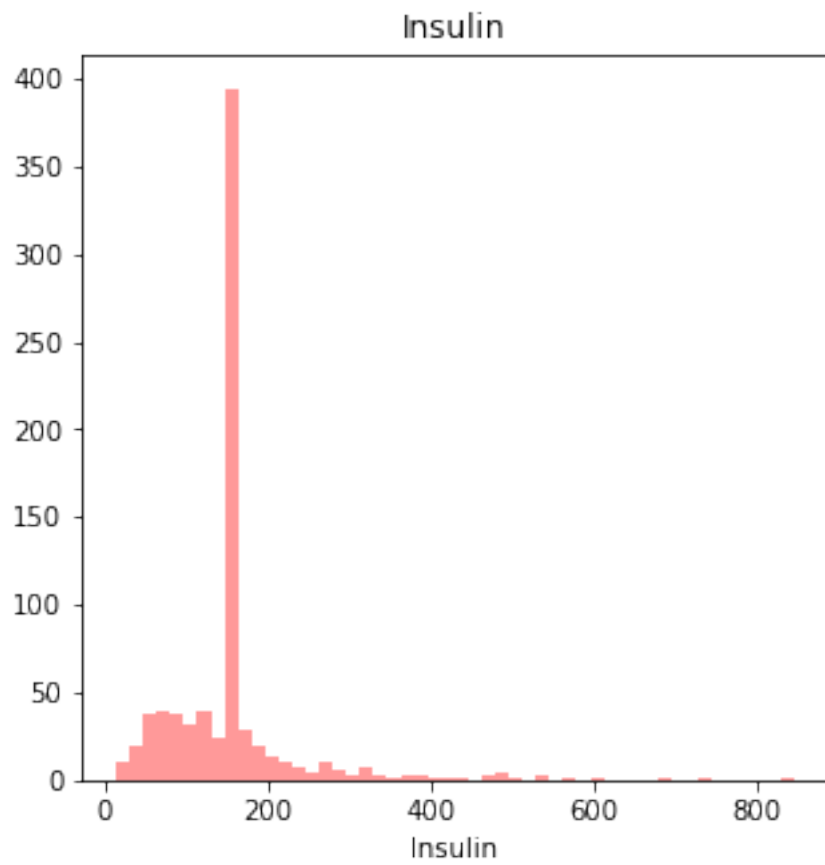
C:\Users\91820\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).


```
warnings.warn(msg, FutureWarning)
```













```
[26]: imputed_data.dtypes
```

```
[26]: Pregnancies      float64
      Glucose        float64
      BloodPressure  float64
      SkinThickness  float64
      Insulin        float64
      BMI            float64
      DiabetesPedigreeFunction float64
      Age            float64
      Outcome        float64
      dtype: object
```

```
[27]: (data.dtypes).value_counts().plot(kind = 'bar')
      plt.show()
```



```
[28]: (imputed_data.Outcome).value_counts().plot(kind = 'bar')  
plt.show()
```



```
[29]: imputed_data.Outcome.value_counts()
```

```
[29]: 0.0    500  
      1.0    268  
      Name: Outcome, dtype: int64
```

```
[30]: round(imputed_data.Outcome.value_counts(normalize = True)*100, 2)
```

```
[30]: 0.0    65.1  
      1.0    34.9  
      Name: Outcome, dtype: float64
```

```
[31]: sns.set()  
      g = sns.pairplot(imputed_data, hue = 'Outcome')  
      g.map_lower(sns.kdeplot)  
      g.map_upper(plt.scatter)  
      g.map_diag(sns.kdeplot)  
      plt.show()
```



```
[32]: round(imputed_data.corr()['Outcome'][:, 3].sort_values(ascending = False)
```

```
[32]: Outcome          1.000
      Glucose          0.494
      BMI             0.314
      Age             0.238
      SkinThickness   0.226
      Pregnancies     0.222
      Insulin         0.214
      DiabetesPedigreeFunction 0.174
      BloodPressure   0.171
      Name: Outcome, dtype: float64
```

```
[33]: def color_negative_red(value):
        """ Colors elements in a dataframe green if positive and red if
        ↪negative. Does not color NaN values."""
        if value < -0.1:
            color = 'red'
        elif value > 0.1:
            color = 'green'
        else:
            color = 'white'
        return 'color: %s' % color
```

```
[34]: round(imputed_data.corr(), 3).style.applymap(color_negative_red)
```

```
[34]: <pandas.io.formats.style.Styler at 0x196ea272ca0>
```

```
[35]: sns.set_style("whitegrid")
corr = data.corr()
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
#kot = corr[corr>=.6]
plt.figure(figsize=(15,10))
sns.heatmap(round(imputed_data.corr(), 3), cmap="coolwarm", vmin=-1,
vmax=1, annot = True, mask = mask, linewidths=1, linecolor='black',
annot_kws={"fontsize":14}).set_title('Correlation Heat Map', fontsize = 20)
plt.grid('on', )
plt.show()
```


C:\Users\91820\AppData\Local\Temp\ipykernel_32068\4172546160.py:3:

DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool_` here.

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
mask = np.zeros_like(corr, dtype=np.bool)
```



```
[36]: # strong correlation seen between:
      #Age & Pregnancies , Glucose & Outcome , BMI & Skin thickness
```

DATA MODELLING

```
[37]: imputed_data.head()
```

```
[37]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6.0    148.0           72.0    35.000000  155.333764  33.6
1           1.0     85.0           66.0    29.000000  155.548223  26.6
2           8.0    183.0           64.0    29.367818  155.374337  23.3
3           1.0     89.0           66.0    23.000000   94.000000  28.1
4           0.0    137.0           40.0    35.000000  168.000000  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627  50.0      1.0
1                0.351  31.0      0.0
2                0.672  32.0      1.0
3                0.167  21.0      0.0
```

4 2.288 33.0 1.0

```
[38]: from sklearn.preprocessing import MinMaxScaler
      scaler = MinMaxScaler()
```

```
[39]: imputed_data_scaled = scaler.fit_transform(imputed_data)
```

```
[40]: imputed_data_scaled = pd.DataFrame(imputed_data_scaled, columns=imputed_data.
      ↪columns)
```

```
[41]: imputed_data_scaled.head()
```

```
[41]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|---|-------------|----------|---------------|---------------|----------|----------|---|
| 0 | 0.352941 | 0.670968 | 0.489796 | 0.304348 | 0.169872 | 0.314928 | |
| 1 | 0.058824 | 0.264516 | 0.428571 | 0.239130 | 0.170130 | 0.171779 | |
| 2 | 0.470588 | 0.896774 | 0.408163 | 0.243128 | 0.169921 | 0.104294 | |
| 3 | 0.058824 | 0.290323 | 0.428571 | 0.173913 | 0.096154 | 0.202454 | |
| 4 | 0.000000 | 0.600000 | 0.163265 | 0.304348 | 0.185096 | 0.509202 | |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|----------|---------|
| 0 | 0.234415 | 0.483333 | 1.0 |
| 1 | 0.116567 | 0.166667 | 0.0 |
| 2 | 0.253629 | 0.183333 | 1.0 |
| 3 | 0.038002 | 0.000000 | 0.0 |
| 4 | 0.943638 | 0.200000 | 1.0 |

```
[42]: from sklearn.linear_model import LogisticRegression
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import classification_report, confusion_matrix
```

```
[43]: imputed_data_scaled.head()
```

```
[43]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|---|-------------|----------|---------------|---------------|----------|----------|---|
| 0 | 0.352941 | 0.670968 | 0.489796 | 0.304348 | 0.169872 | 0.314928 | |
| 1 | 0.058824 | 0.264516 | 0.428571 | 0.239130 | 0.170130 | 0.171779 | |
| 2 | 0.470588 | 0.896774 | 0.408163 | 0.243128 | 0.169921 | 0.104294 | |
| 3 | 0.058824 | 0.290323 | 0.428571 | 0.173913 | 0.096154 | 0.202454 | |
| 4 | 0.000000 | 0.600000 | 0.163265 | 0.304348 | 0.185096 | 0.509202 | |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|----------|---------|
| 0 | 0.234415 | 0.483333 | 1.0 |
| 1 | 0.116567 | 0.166667 | 0.0 |
| 2 | 0.253629 | 0.183333 | 1.0 |
| 3 | 0.038002 | 0.000000 | 0.0 |
| 4 | 0.943638 | 0.200000 | 1.0 |

```
[44]: y = imputed_data_scaled['Outcome']
      x = imputed_data_scaled.drop('Outcome', axis = 1)
```

```
[45]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size= 0.2,
↳stratify = y)
```

```
[46]: x_train.head()
```

```
[46]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin      BMI  \
542      0.588235  0.296774      0.622449      0.271739  0.170130  0.341513
646      0.058824  0.793548      0.510204      0.108696  0.156250  0.106339
525      0.176471  0.277419      0.367347      0.119565  0.170355  0.073620
386      0.294118  0.464516      0.510204      0.239130  0.170130  0.288344
157      0.058824  0.419355      0.326531      0.152174  0.145433  0.143149

      DiabetesPedigreeFunction      Age
542              0.318958  0.583333
646              0.157558  0.200000
525              0.156277  0.000000
386              0.248506  0.233333
157              0.322374  0.033333
```

```
[47]: y_train.head()
```

```
[47]: 542      1.0
      646      1.0
      525      0.0
      386      1.0
      157      0.0
      Name: Outcome, dtype: float64
```

```
[48]: lr=LogisticRegression()
```

```
[49]: lr.fit(x_train, y_train)
```

```
[49]: LogisticRegression()
```

```
[50]: pred = lr.predict(x_test)
```

```
[51]: cnf_matrix = confusion_matrix(y_test, pred)
```

```
[52]: cnf_matrix
```

```
[52]: array([[89, 11],
      [26, 28]], dtype=int64)
```



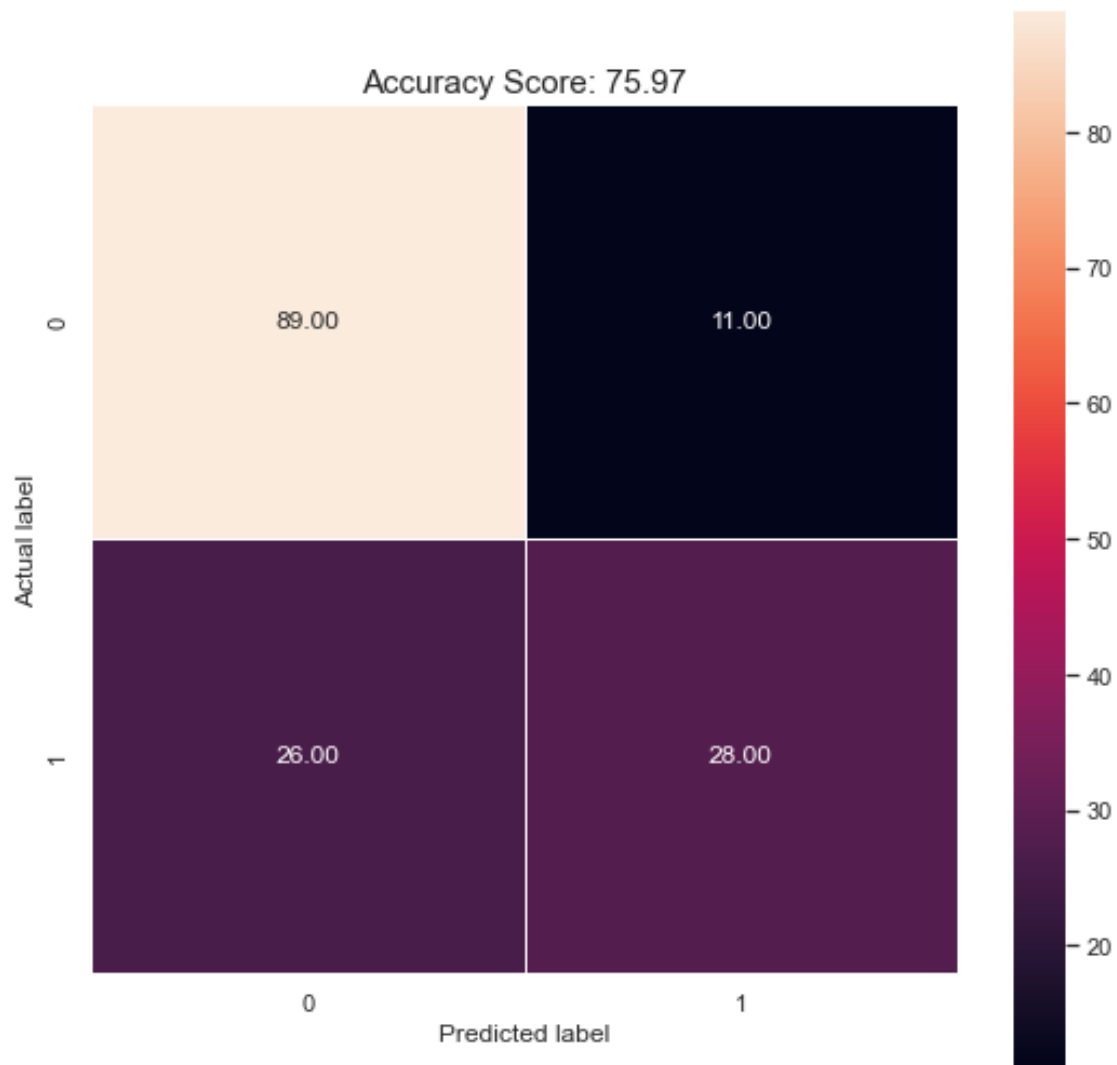
```
[53]: def sens_spec(cnf_matrix):
    total_cm = sum(sum(cnf_matrix))
    accuracy_clf = (cnf_matrix[0,0] + cnf_matrix[1,1]) / total_cm
    sensitivity_clf = cnf_matrix[0,0] / (cnf_matrix[0, 0] + cnf_matrix[0, 1])
    specificity_clf = cnf_matrix[1,1] / (cnf_matrix[1, 0] + cnf_matrix[1, 1])

    #print('accuracy of {} is {}'.format(accuracy))
    return('Accuracy: {}'.format(round(accuracy_clf, 2)), 'Sensitivity: {}'.
    ↪format(round(sensitivity_clf, 2)), 'Specificity: {}'.
    ↪format(round(specificity_clf, 2)))
```

```
[54]: # Use score method to get accuracy of model
score = lr.score(x_test, y_test)
print(score)
```

0.7597402597402597

```
[55]: plt.figure(figsize=(9,9))
sns.heatmap(cnf_matrix, annot=True, fmt=".2f", linewidths=.5, square = True);
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Accuracy Score: {}'.format(round(score*100, 2
))
plt.title(all_sample_title, size = 15);
```



```
[56]: print(sens_spec(cnf_matrix))
      print()
      print(classification_report(y_test, pred))
```

('Accuracy: 0.76', 'Sensitivity: 0.89', 'Specificity: 0.52')

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.77 | 0.89 | 0.83 | 100 |
| 1.0 | 0.72 | 0.52 | 0.60 | 54 |
| accuracy | | | 0.76 | 154 |
| macro avg | 0.75 | 0.70 | 0.72 | 154 |
| weighted avg | 0.75 | 0.76 | 0.75 | 154 |

```
[57]: from imblearn.over_sampling import SMOTE
```

```
[58]: os = SMOTE(random_state=0)
```

```
[59]: columns = x_train.columns
```

```
[60]: os_data_X,os_data_y=os.fit_resample(x, y)
os_data_X = pd.DataFrame(data=os_data_X,columns=columns )
os_data_y= pd.DataFrame(data=os_data_y,columns=['Outcome'])
# we can Check the numbers of our data
print("length of oversampled data is ",len(os_data_X))
print("Number of NEGATIVE in oversampled_
↳data",len(os_data_y[os_data_y['Outcome']==0]))
print("Number of POSITIVE",len(os_data_y[os_data_y['Outcome']==1]))
print("Proportion of NEGATIVE data in oversampled data is_
↳",len(os_data_y[os_data_y['Outcome']==0])/len(os_data_X))
print("Proportion of POSITIVE data in oversampled data is_
↳",len(os_data_y[os_data_y['Outcome']==1])/len(os_data_X))
```

length of oversampled data is 1000

Number of NEGATIVE in oversampled data 500

Number of POSITIVE 500

Proportion of NEGATIVE data in oversampled data is 0.5

Proportion of POSITIVE data in oversampled data is 0.5

```
[61]: X_train, X_test, y_train, y_test = train_test_split(os_data_X,os_data_y,
↳test_size=0.2, random_state=0)
```

```
[62]: from sklearn.neural_network import MLPClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.gaussian_process import GaussianProcessClassifier
from sklearn.gaussian_process.kernels import RBF
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
```

```
[65]: import scikitplot as skplt
```

```
↳
-----
ModuleNotFoundError                                Traceback (most recent call_
↳last)
```

```
~\AppData\Local\Temp\ipykernel_32068\105772748.py in <module>
----> 1 import scikitplot as skplt
```

```
ModuleNotFoundError: No module named 'scikitplot'
```

```
[66]: names = ["Nearest Neighbors", "Logistic Regression", "Linear SVM",
               "RBF SVM", "Gaussian Process",
               "Decision Tree", "Random Forest", "Neural Net", "AdaBoost"
               ,
               "Naive Bayes", "QDA"]
       classifiers = [
           KNeighborsClassifier(3),
           LogisticRegression(),
           SVC(kernel="linear", C=0.025, probability=True),
           SVC(gamma=2, C=1, probability=True),
           GaussianProcessClassifier(1.0 * RBF(1.0)),
           DecisionTreeClassifier(max_depth=5),
           RandomForestClassifier(max_depth=5, n_estimators=10,
               ↳max_features=1), MLPClassifier(alpha=1, max_iter=1000),
           AdaBoostClassifier(), GaussianNB(), QuadraticDiscriminantAnalysis()]
```

```
[82]: # iterate over classifiers
       for name, clf in zip(names, classifiers):
           #ax = plt.subplot(len(datasets), len(classifiers) + 1, i)
           print("classifier:", name)
           clf.fit(X_train, y_train)
           score = clf.score(X_test, y_test)
           pred = clf.predict(X_test)
           prob = clf.predict_proba(X_test)
           print(round(score*100, 2))
           print(sens_spec(cnf_matrix))
           print()
           cnf_matrix = confusion_matrix(y_test, pred)
           plt.figure(figsize=(9,9))
           sns.heatmap(cnf_matrix, annot=True, fmt=".2f", linewidths=.5, square = True);
           plt.ylabel('Actual label');
           plt.xlabel('Predicted label');
           #all_sample_title = 'Accuracy Score: {0}'.format(round(score*100, 2))
           all_sample_title = 'Classifier: {}, Accuracy Score: {}'.format(name,
               ↳round(score*100, 2))
           plt.title(all_sample_title, size = 15);
           #print()
           print(classification_report(y_test, pred))
           print()
```

```

skplt.metrics.plot_roc_curve(y_test, prob, figsize = (15, 10), title = 'ROC_
↳Curve for: {}'.format(name))
plt.show()
print('-----')
print()

```

```

classifier: Nearest Neighbors
classifier: Logistic Regression
classifier: Linear SVM
classifier: RBF SVM
classifier: Gaussian Process
classifier: Decision Tree
classifier: Random Forest
classifier: Neural Net
classifier: AdaBoost
classifier: Naive Bayes
classifier: QDA
77.5
('Accuracy: 0.78', 'Sensitivity: 0.83', 'Specificity: 0.72')

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.76 | 0.83 | 0.79 | 105 |
| 1.0 | 0.79 | 0.72 | 0.75 | 95 |
| accuracy | | | 0.78 | 200 |
| macro avg | 0.78 | 0.77 | 0.77 | 200 |
| weighted avg | 0.78 | 0.78 | 0.77 | 200 |

```

/usr/local/lib/python3.7/site-packages/sklearn/utils/validation.py:63:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().

```

```

    return f(*args, **kwargs)

```

```

↳-----

```

```

TypeError                                Traceback (most recent call_
↳last)

```

```

<ipython-input-86-a16a05fc551c> in <module>
    21 print(classification_report(y_test, pred))
    22 print()

```

```

---> 23 sk.metrics.plot_roc_curve(y_test, prob, figsize = (15, 10), title =
↳ 'ROC Curve for: {}'.format(name))
    24 plt.show()
    25
↳ print('-----


```

```

    /usr/local/lib/python3.7/site-packages/sklearn/utils/validation.py in
↳ inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
---> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

```

TypeError: plot_roc_curve() missing 1 required positional argument: 'y'



```
[67]: #import random
np.random.seed(1000)
randomlist = []
for i in range(0,10):
    n = np.random.randint(1,len(X_test))
    randomlist.append(n)
    print(randomlist)
```

```
[180]
[180, 88]
[180, 88, 72]
[180, 88, 72, 193]
[180, 88, 72, 193, 95]
[180, 88, 72, 193, 95, 93]
[180, 88, 72, 193, 95, 93, 2]
```

```
[180, 88, 72, 193, 95, 93, 2, 190]
[180, 88, 72, 193, 95, 93, 2, 190, 129]
[180, 88, 72, 193, 95, 93, 2, 190, 129, 90]
```

```
[68]: list(X_test.iloc[180])
```

```
[68]: [0.058823529411764705,
       0.3161290322580645,
       0.4693877551020407,
       0.2608695652173913,
       0.1701499109153931,
       0.24948875255623731,
       0.10119555935098205,
       0.033333333333333326]
```

```
[69]: y_test.iloc[180]['Outcome']
```

```
[69]: 0.0
```

```
[70]: pre_out = []
       out = []
       for i in randomlist:
           data_in = [list(X_test.iloc[i])]
           data_in = np.around(data_in, 2)
           pre_data_out = lr.predict(data_in)
           data_out = y_test.iloc[i]['Outcome']
           mylist = [i, data_in, pre_data_out, data_out]
           print(*mylist, sep='\n')
           print('-----')
       pre_out.append(pre_data_out)
       out.append(data_out)
```

```
180
[[0.06 0.32 0.47 0.26 0.17 0.25 0.1  0.03]]
[0.]
0.0
-----
88
[[0.12 0.24 0.49 0.09 0.07 0.24 0.2  0.07]]
[0.]
0.0
-----
72
[[0.65 0.49 0.57 0.33 0.16 0.49 0.3  0.45]]
[1.]
1.0
-----
193
```



```
[[0.29 0.46 0.51 0.24 0.17 0.15 0.05 0.15]]
[0.]
0.0
```

```
-----
95
[[0.    0.59 0.45 0.38 0.28 0.49 0.12 0.05]]
[0.]
1.0
```

```
-----
93
[[0.12 0.41 0.39 0.03 0.32 0.15 0.34 0.02]]
[0.]
0.0
```

```
-----
2
[[0.82 0.36 0.55 0.2  0.2  0.38 0.14 0.42]]
[0.]
1.0
```

```
-----
190
[[0.52 0.46 0.57 0.27 0.19 0.32 0.08 0.26]]
[0.]
1.0
```

```
-----
129
[[0.    0.7  0.59 0.35 0.31 0.48 0.08 0.1 ]]
[1.]
0.0
```

```
-----
90
[[0.    0.3  0.45 0.27 0.24 0.44 0.13 0.07]]
[0.]
0.0
```

```
-----
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but LogisticRegression was fitted with
feature names
```

```
    warnings.warn(
```

```
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but LogisticRegression was fitted with
feature names
```

```
    warnings.warn(
```

```
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but LogisticRegression was fitted with
feature names
```

```
    warnings.warn(
```

```
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
```

```

does not have valid feature names, but LogisticRegression was fitted with
feature names
    warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but LogisticRegression was fitted with
feature names
    warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but LogisticRegression was fitted with
feature names
    warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but LogisticRegression was fitted with
feature names
    warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but LogisticRegression was fitted with
feature names
    warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but LogisticRegression was fitted with
feature names
    warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but LogisticRegression was fitted with
feature names
    warnings.warn(

```

```

[71]: svc = SVC(gamma=2, C=1, probability=True)
      svc.fit(X_train, y_train)

```

```

C:\Users\91820\anaconda3\lib\site-packages\sklearn\utils\validation.py:993:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)

```

```

[71]: SVC(C=1, gamma=2, probability=True)

```

```

[72]: pre_out = []
      out = []
      for i in randomlist:
          data_in = [list(X_test.iloc[i])]
          data_in = np.around(data_in, 2)
          pre_data_out = svc.predict(data_in)
          data_out = y_test.iloc[i]['Outcome']
          mylist = [i, data_in, pre_data_out, data_out]
          print(*mylist, sep='\n')

```

```

print('-----')
pre_out.append(pre_data_out)
out.append(data_out)

```

```

180
[[0.06 0.32 0.47 0.26 0.17 0.25 0.1 0.03]]
[0.]
0.0
-----
88
[[0.12 0.24 0.49 0.09 0.07 0.24 0.2 0.07]]
[0.]
0.0
-----
72
[[0.65 0.49 0.57 0.33 0.16 0.49 0.3 0.45]]
[1.]
1.0
-----
193
[[0.29 0.46 0.51 0.24 0.17 0.15 0.05 0.15]]
[0.]
0.0
-----
95
[[0.    0.59 0.45 0.38 0.28 0.49 0.12 0.05]]
[1.]
1.0
-----
93
[[0.12 0.41 0.39 0.03 0.32 0.15 0.34 0.02]]
[0.]
0.0
-----
2
[[0.82 0.36 0.55 0.2 0.2 0.38 0.14 0.42]]
[1.]
1.0
-----
190
[[0.52 0.46 0.57 0.27 0.19 0.32 0.08 0.26]]
[1.]
1.0
-----
129
[[0.    0.7 0.59 0.35 0.31 0.48 0.08 0.1 ]]
[1.]

```

```

0.0
-----
90
[[0.    0.3  0.45 0.27 0.24 0.44 0.13 0.07]]
[0.]
0.0
-----

```

```

C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(
C:\Users\91820\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X
does not have valid feature names, but SVC was fitted with feature names
  warnings.warn(

```

1 Our Basic Logistic Regression algorithm is

predicting 6 / 10 inputs correctly, whereas the best performing SVC algorithm is predicting 9 / 10 inputs correctly

[]:

[]:

[]: